

Data Driven Decision Making: Sampling Distributions & Estimation

σ = std dev of popⁿ (within)
 S : std dev of sample (within)
 $\frac{\sigma}{\sqrt{n}}$ or $\frac{S}{\sqrt{n}}$ std error \Rightarrow std dev b/w \bar{x} values (sample averages)

GSBA 545, Fall 2021

Professor Dawn Porter

- Sampling
 - Expected Value & Standard Error of Sample Means
- Application - Probability of Averages
- Confidence Intervals
 - Large Sample, Means
 - Small Sample, Means
 - Population Proportions
- Sample Size Determination
 - Means
 - Proportions

Estimating Parameters

Parameter: a characteristic of the population

Statistic: an observed characteristic of a sample

Name	Sample Statistic	Population Parameter
Mean	\bar{x}	μ
Standard deviation	s	σ
Correlation	r	ρ
Proportion	\hat{p}	p
Regression intercept	b_0	β_0
Regression slope	b_i	β_i

Sampling Variation:

- The variability in the value of a statistic from sample to sample.
- The price we pay for working with a sample rather than the population.

Assume two separate samples are taken to estimate the population average income in LA:

$$n_1 = 10, \text{ sample average} = \$45,000, \quad n_2 = 1000, \text{ sample average} = \$45,000$$

- Which result would you trust more? Why?
- If the experiment were repeated, which result would vary more? Why?

Sampling Distributions

If a random sample of size n is taken from a (normal) population with mean μ and standard deviation σ , then the sampling distribution of the sample mean \bar{x} has

mean:
standard error:

$$E(\bar{x}) = \mu_{\bar{x}} = \mu$$
$$SE(\bar{x}) = \sigma_{\bar{x}} = \frac{\sigma}{\sqrt{n}}$$



$n \geq 30$ for normality.

1. The best guess of μ is \bar{x} and the reverse.
2. As n increases, estimates become more reliable.
3. Regardless of underlying population distribution, possible sample averages should be normally distributed (if $n > 30$).

*CV is by how much population averages differ, std is how the observations vary.
Watch 09/22 video.*


Sample Means

$n = 3$ stocks in each sample					
	Sample	Stock Returns			Mean
1	A,B,C	-36	-15	3	-16




Sample Size Effect


(a) The population of individual mileages



(b) The sampling distribution of the sample mean \bar{x} when $n = 5$

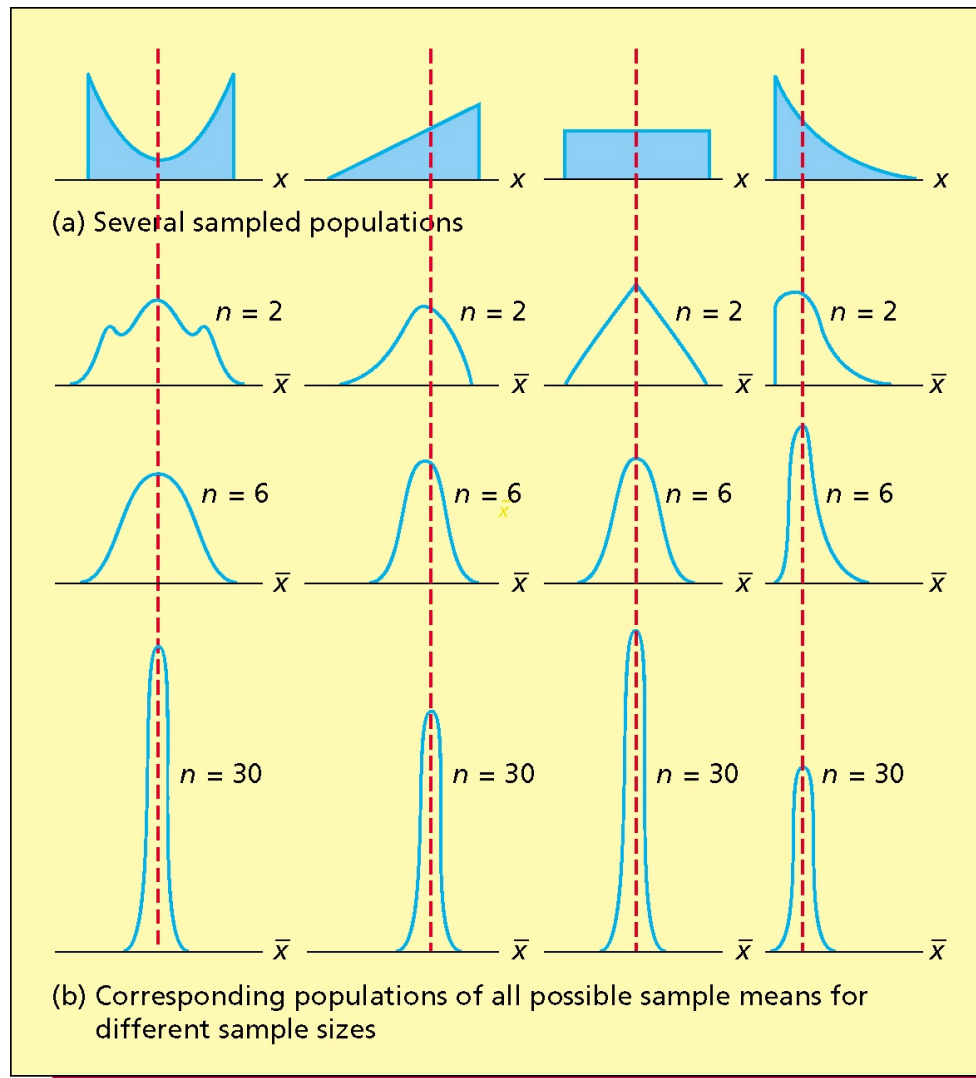


(c) The sampling distribution of the sample mean \bar{x} when $n = 49$



Sample Size Effect

The larger the sample size, the more nearly normally distributed is the population of all possible sample means, even if the original distribution was far from normal.



Why is this concept important and where is it used?

- If population parameters are known, we can calculate the probability of certain sample averages occurring.
- In Statistical Inference, we use sample results to test hypotheses.
- In Regression Analysis, sample results (including standard errors!) are needed to draw intelligent conclusions.

Probability of Averages

Assume the prices of meals at a restaurant are normally distributed with an average of \$35 and a standard deviation of \$5. I have \$33 in my pocket.

1. What is the probability I have enough money for my meal?

$$P(x < 33) = P\left(z \leq \frac{33-35}{5}\right) = P(z \leq -0.4) = 0.3446 \Rightarrow 34.46\%$$

2. I'm out to dinner with friends and we're splitting the check evenly by 10.

a) What is the expected price per meal now?

b) What is the probability I have enough money now?

c) Why did the probability change? \Rightarrow since we avg over 10, std error is less & values converge towards \$35.

$$P(\bar{x} < 33) = P\left(z \leq \frac{33-35}{5/\sqrt{10}}\right) = P(z \leq -1.26) = 0.1038 \Rightarrow 10.38\%$$

Sample estimates are generally point estimates, but these are approximations.


Therefore, we need to have an understanding of the data to create a **margin of error** and create an interval around a point estimate.

The margin of error is computed using either:

- the population standard deviation σ , or
- the sample standard deviation s
- σ is rarely known exactly, but often a good estimate can be obtained based on historical data or other information; in that case, we can assume s is known


Margin of Error

There is a 0.95 probability that the value of a sample mean will provide a margin of error of $1.96 \times \frac{\sigma}{\sqrt{n}}$ or less.




Margin of Error

There is a $(1 - \alpha)$ probability that the value of a sample mean will provide a margin of error of $z_{\alpha/2} \frac{\sigma}{\sqrt{n}}$ or less.



Margin of Error

There is a $(1 - \alpha)$ probability μ will be within a margin of error of $z_{\alpha/2} \frac{\sigma}{\sqrt{n}}$ or less from a sample mean.



If σ is known, or if $n \geq 30$, we can be $100(1 - \alpha)\%$ confident that μ falls between:

$$\bar{x} \pm z_{\alpha/2} \frac{\sigma}{\sqrt{n}} \quad \text{or} \quad \bar{x} \pm z_{\alpha/2} \frac{s}{\sqrt{n}}$$

where:

\bar{x} is the sample mean

$(1 - \alpha)$ is the confidence coefficient

$z_{\alpha/2}$ is z value for an area of $\alpha/2$ in the upper tail

σ (s) is the population (sample) standard deviation

n is the sample size

Example: Discount Sounds

Discount Sounds has several retail outlets throughout the United States and the firm is evaluating a potential location for a new outlet, based in part on the mean annual income of the individuals in the marketing area of the new location.

A sample of size $n = 36$ individuals was taken; the sample mean income is \$41,100. The population is not believed to be highly skewed. The sample standard deviation is calculated to be \$4,500, and the confidence level to be used is 95%.

Create a 95% Confidence Interval for the *population mean income*.

$$\bar{x} \pm z_{\frac{\alpha}{2}} \frac{s}{\sqrt{n}} \Rightarrow 41,000 \pm 1.96 \frac{4,500}{\sqrt{36}} \Rightarrow 41,000 \pm \underbrace{1470}_{\text{margin of error}}$$


Example: Discount Sounds

Various confidence level intervals

Confidence level		
90%		
95%		
99%		

- If $n < 30$, an adjustment needs to be made to account for a higher level of uncertainty.
- Instead of using the **Normal Distribution** and Z values, we rely on the **t -Distribution**.
- Using the **t -Distribution** is more conservative because its shape is more spread out than the Normal; as the sample size increases, though, the two curves will converge.
- To use this distribution, we need the degrees of freedom, or df . In these interval calculations, $df = n - 1$.

Effect of df on t -distribution



When $df > 100$, the z values are good approximations to the t value.

Standard normal Z values can be found on the infinite degrees (∞) row of the t -distribution table.

If $n < 30$, an adjustment needs to be made to account for a higher level of uncertainty. The new interval for μ , at a $100(1 - \alpha)\%$ confidence level, is:

$$\bar{x} \pm t_{\alpha/2} \frac{s}{\sqrt{n}}$$

where:

- \bar{x} is the sample mean
- $(1 - \alpha)$ is the confidence coefficient
- $t_{\alpha/2}$ is t value for an area of $\alpha/2$ in the upper tail
- s is the sample standard deviation
- n is the sample size

Example: Apartment Rents

A reporter for a student newspaper is writing an article on the cost of off-campus housing. A sample of 16 one-bedroom apartments within a half-mile of campus resulted in a sample mean of \$750 per month and a sample standard deviation of \$55.

Create a 95% confidence interval estimate of the mean rent per month for the population of one-bedroom apartments within a half-mile of campus. We will assume this population to be normally distributed.

If the goal is to estimate the *population proportion*, p , we can create a $100(1 - \alpha)\%$ confident that p falls between:

$$\hat{p} \pm z_{\alpha/2} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

where:

\hat{p}	is the sample proportion
$(1 - \alpha)$	is the confidence coefficient
$z_{\alpha/2}$	is z value with $\alpha/2$ in the upper tail
$\hat{p}(1 - \hat{p})$	is an estimate of the std deviation
n	is the sample size

Example: Clinical Trial Results

A new drug called Phe-Myecin is being marketed by a company and to get through the clinical trial requirements, the company must forecast what percentage of the population of drug-takers would experience negative side effects.

A sample of $n = 200$ patients was taken and 35 people experienced negative side effects. What is a 95% range for what the true population proportion experiencing negative side effects will be?

To calculate a reasonable sample size to be taken to obtain an estimate of μ , the margin of error needs to be specified.

Margin of Error:

$$E = z_{\alpha/2} \frac{\sigma}{\sqrt{n}}$$

Then the necessary sample size is:

$$n = \left(\frac{z_{\alpha/2} \sigma}{E} \right)^2$$

Note that a value for σ is necessary here. If σ is unknown, either:

1. Use a reasonable estimate of σ computed in a previous study.
2. Use a pilot study to select a preliminary sample to estimate and use s from there.

Example: Discount Sounds

Recall that Discount Sounds is evaluating a potential location for a new retail outlet, based in part, on the mean annual income of the individuals in the marketing area of the new location.

Suppose that Discount Sounds' management team wants to create a 95% confidence interval with a margin of error of \$500 or less. How large a sample size is needed for this? (Recall from earlier that $s = 4500$.)

Sample Sizes: Proportions

To calculate a reasonable sample size to be taken to obtain an estimate of p , the margin of error needs to be specified.

Margin of Error:
$$E = z_{\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

Then the necessary sample size is:
$$n = p(1 - p) \left(\frac{z_{\alpha/2}}{E} \right)^2$$

Note that a value for p is necessary here, which isn't known yet. Options are:

1. Use a reasonable estimate of \hat{p} computed in a previous study.
2. Use a pilot study to select a preliminary sample and use that \hat{p} .
3. Otherwise, use $p = 0.50$ as a conservative starting point.

Example: Clinical Trial Results

Assume the interval result from the clinical trial of Phe-Myecin is too wide to get approval. Now the company needs to report a 95% confidence interval with a margin of error of only $\pm 4\%$ at the most. How large a sample size is needed for this if we do use our prior estimate of $\hat{p} = 0.175$?

Example: Clinical Trial Results (limited information)

Now assume the company hasn't done any study yet and is trying to work out a reasonable sample size for the trial. They are told they need to report a 95% confidence interval with a margin of error of only $\pm 4\%$ at the most. How large a sample size is needed for this if we do not use our prior estimate of \hat{p} ?